

# Toxic Comment Classification

## Introduction

Discussing things you care about can be difficult. The threat of abuse and harassment online means that many people stop expressing themselves and give up on seeking different opinions. Platforms struggle to effectively facilitate conversations, leading many communities to limit or completely shut down user comments.

One area of focus is the study of negative online behaviors, like toxic comments (i.e. comments that are rude, disrespectful or otherwise likely to make someone leave a discussion).

Here, we would try to build a multi-headed model that's capable of detecting different types of toxicity like threats, obscenity, insults, and identity-based hate.

ref: <https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge/>  
(<https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge/>)

## Data

We'll be using a dataset of comments from Wikipedia's talk page edits. Improvements to the current model will hopefully help online discussion become more productive and respectful. The corpus contains 63M comments from discussions relating to user pages and articles dating from 2004-2015.

*Disclaimer: the dataset for this competition contains text that may be considered profane, vulgar, or offensive.*

We are provided with a large number of Wikipedia comments which have been labeled by human raters for toxic behavior. The types of toxicity are:

- toxic
- severe\_toxic
- obscene
- threat
- insult
- identity\_hate

The tagging was done via crowdsourcing which means that the dataset was rated by different people and the tagging might not be 100% accurate too.

The source paper also contains more interesting details about the dataset creation. (link to source paper: <https://arxiv.org/pdf/1610.08914.pdf> (<https://arxiv.org/pdf/1610.08914.pdf>))

```
In [1]: import pandas as pd
import numpy as np

from sklearn.feature_extraction.text import TfidfVectorizer
```

train.csv is the training set containing comments with their binary labels for each toxicity category

## 1. Reading Data

```
In [2]: data=pd.read_csv("train.csv")
data.head()
```

Out[2]:

	id	comment_text	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	0000997932d777bf	Explanation\nWhy the edits made under my usern...	0	0	0	0	0	0
1	000103f0d9cfb60f	D'aww! He matches this background colour I'm s...	0	0	0	0	0	0
2	000113f07ec002fd	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0
3	0001b41b1c6bb37e	"\nMore\nI can't make anyreal suggestions on ...	0	0	0	0	0	0
4	0001d958c54c6e35	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0

```
In [3]: test_data=pd.read_csv("test.csv")
test_data.head(3)
```

Out[3]:

	id	comment_text
0	00001cee341fdb12	Yo bitch Ja Rule is more succesful then you'll...
1	0000247867823ef7	== From RfC == \n\n The title is fine as it is...
2	00013b17ad220c46	"\n\n == Sources == \n\n * Zawe Ashton on Lap...

## 2. Understanding Data

### sample comment text

```
In [4]: data.loc[2, 'comment_text']
```

Out[4]: "Hey man, I'm really not trying to edit war. It's just that this guy is constantly removing relevant information and talking to me through edits instead of my talk page. He seems to care more about the formatting than the actual info."

```
In [5]: data.loc[33, 'comment_text']
```

Out[5]: "I was able to post the above list so quickly because I already had it in a text file in my hard drive I've been meaning to get around to updating the sound list for some time now. \nAs far as generating interest I've spent four years trying to drum up more interest in freely licensed full length classical music. Unfortunately, my attempts failed - I'm still effectively the only one who does it. The classical music wikiproject was not interested, (Wikipedia\_talk:WikiProject\_Classical\_music/Archive\_5#Need\_help.21Wikipedia\_talk:WikiProject\_Music/Archive\_3#I\_could\_use\_some\_helpWikipedia\_talk:WikiProject\_Music/Archive\_2#Raulbot.2C\_and\_the\_music\_list) So I really had given up trying to interest others. \nThe sound list was featured on digg a while back - [http://digg.com/music/Wikipedia\\_has\\_free\\_classical\\_music\\_downloads](http://digg.com/music/Wikipedia_has_free_classical_music_downloads) ([http://digg.com/music/Wikipedia\\_has\\_free\\_classical\\_music\\_downloads](http://digg.com/music/Wikipedia_has_free_classical_music_downloads)) . It got 1600 diggs, which is IMO very impressive."

## Comment text length analysis

```
In [6]: data['comment_text'].str.len().describe()
```

```
Out[6]: count      159571.000000
mean         394.073221
std          590.720282
min           6.000000
25%          96.000000
50%         205.000000
75%         435.000000
max        5000.000000
Name: comment_text, dtype: float64
```

We can see the lengths of comment varies a lot as largest comment has 5000 words while smallest comment has length of only 6 words.

Also, the mean and median of length differs a lot as well (mean of 394 words vs. median of 205 words).

```
In [7]: labels=data.columns.tolist()
labels.remove('id')
labels.remove('comment_text')
print(labels)
```

```
['toxic', 'severe_toxic', 'obscene', 'threat', 'insult', 'identity_hate']
```

```
In [8]: #Analysis: tags per comments
data['label_num']= data[labels].sum(axis=1)
data.head()
```

```
Out[8]:
```

	id	comment_text	toxic	severe_toxic	obscene	threat	insult	identity_hate	label_num
0	0000997932d777bf	Explanation\nWhy the edits made under my usern...	0	0	0	0	0	0	
1	000103f0d9cfb60f	D'aww! He matches this background colour I'm s...	0	0	0	0	0	0	
2	000113f07ec002fd	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0	
3	0001b41b1c6bb37e	"\nMore\nI can't make any real suggestions on ...	0	0	0	0	0	0	
4	0001d958c54c6e35	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0	

```
In [9]: data['label_num'].value_counts()
```

```
Out[9]: 0      143346
        1       6360
        3       4209
        2       3480
        4       1760
        5        385
        6         31
        Name: label_num, dtype: int64
```

```
In [10]: print("Out of total", data['label_num'].value_counts().sum(), "comments,", data['
         "are non-toxic.")
```

Out of total 159571 comments, 143346 are non-toxic.

```
In [11]: import seaborn as sns
         import matplotlib.pyplot as plot
```

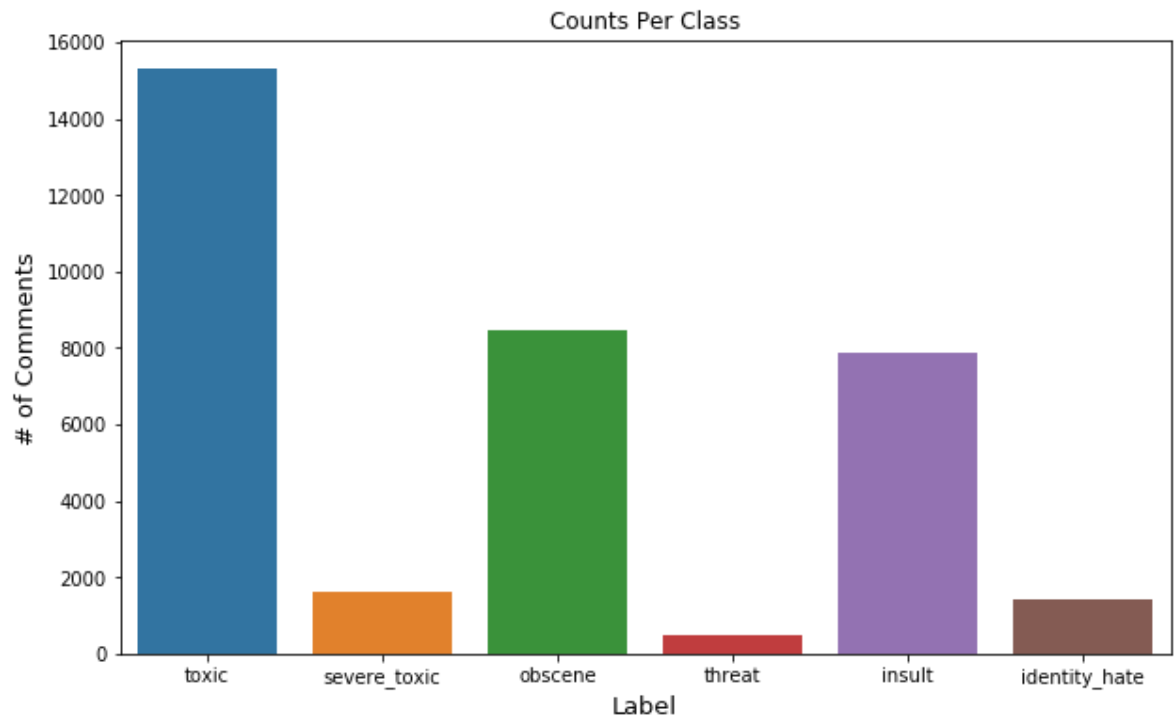
No. of comments per each toxic category

```
In [12]: data[labels].sum(axis=0)
```

```
Out[12]: toxic          15294
         severe_toxic    1595
         obscene         8449
         threat          478
         insult          7877
         identity_hate    1405
         dtype: int64
```

```
In [13]: # plotting the label counts
plot.figure(figsize=(10,6))
ax = sns.barplot(data[labels].sum(axis=0).index, data[labels].sum(axis=0).values)
plot.title("Counts Per Class")
plot.ylabel('# of Comments', fontsize=13)
plot.xlabel('Label', fontsize=13)
```

Out[13]: Text(0.5,0,'Label')



**Sample comments for each toxic category**

```
In [14]: # Analysing toxic text for each category
# Disclaimer: text may contain profane, vulgar, or offensive language.
for lab in labels:
    print("label: ", lab)
    print("Comment text sample:")
    print(data[data[lab]==1].iloc[3, 1])
    print("-----")
```

```
label: toxic
Comment text sample:
You are gay or antisemitian?
```

```
Archangel White Tiger
```

```
Meow! Greetingshhh!
```

```
Uh, there are two ways, why you do erased my comment about WW2, that holocaust was brutally slaying of Jews and not gays/Gypsies/Slavs/anyone...
```

```
1 - If you are anti-semitian, then shave your head bald and go to the skinhead meetings!
```

```
2 - If you doubt words of the Bible, that homosexuality is a deadly sin, make a pentagram tattoo on your forehead go to the satanic masses with your gay pals!
```

```
3 - First and last warning, you fucking gay - I won't appreciate if any more nazis would write in my page! I don't wish to talk to you anymore!
```

```
Beware of the Dark Side!
```

```
-----
label: severe_toxic
Comment text sample:
Hi
```

```
Im a fucking bitch.
```

```
50.180.208.181
```

```
-----
label: obscene
Comment text sample:
GET FUCKED UP. GET FUCKEED UP. GOT A DRINK THAT YOU CANT PUT DOWN???/ GET FUCK UP GET FUCKED UP. I'M FUCKED UP RIGHT NOW!
```

```
-----
label: threat
Comment text sample:
Fuck you, Smith. Please have me notified when you die. I want to dance on your grave.
```

```
-----
label: insult
Comment text sample:
Stupid piece of shit stop deleting my stuff asshole go die and fall in a hole go to hell!
```

```
-----
label: identity_hate
Comment text sample:
Kill all niggers.
```

```
I have heard, that others have said this.. should this be included? That racists sometimes say these.
```

## 2.1 Text data Cleaning

```
In [15]: # text cleaning
# https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\'re", " are", phrase)
    phrase = re.sub(r"\s", " is", phrase)
    phrase = re.sub(r"\d", " would", phrase)
    phrase = re.sub(r"\ll", " will", phrase)
    phrase = re.sub(r"\t", " not", phrase)
    phrase = re.sub(r"\ve", " have", phrase)
    phrase = re.sub(r"\m", " am", phrase)
    return phrase
```

```
In [16]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "
            'you'll", 'you'd', 'your', 'yours', 'yourself', 'yourselves', 'he', '
            'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itsel
            'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that
            'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has
            'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because'
            'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'th
            'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off
            'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all'
            'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than',
            's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've
            've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "di
            'hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma',
            'mustn't', 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn'
            'won', "won't", 'wouldn', "wouldn't"]
```

```
In [17]: from nltk.stem.wordnet import WordNetLemmatizer
lemma = WordNetLemmatizer()
```


```
In [18]: # Combining all the above students
from tqdm import tqdm as tqdm
def preprocess_text(text_data):
    preprocessed_text = []
    # tqdm is for printing the status bar
    for sentence in tqdm(text_data):
        sent = decontracted(sentence)
        sent = sent.replace('\r', ' ')
        sent = sent.replace('\n', ' ')
        sent = sent.replace('\\"', ' ')
        sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
        words = [lemma.lemmatize(word, "v") for word in sent.split()]
        sent = ' '.join(e for e in words if e.lower() not in stopwords)
        preprocessed_text.append(sent.lower().strip())
    return preprocessed_text
```





```
In [22]: imp_dfs = get_topn_tfidf_feat_byClass(X_unigrams, data, feature_names, labels, to
plot.figure(figsize=(15,10))

for i, label in enumerate(labels):
    plot.subplot(3, 2, i + 1)
    sns.barplot(imp_dfs[label].word_feature[:10], imp_dfs[label].tfidf_value[:10])
    plot.title("Important words for the class:{}".format(label))
    plot.tight_layout()
```



**most common words for each category**



```
comment_words = ''
for string in data[data.severe_toxic==1].cleared_text1.values:
    comment_words += string + ' '

wordcloud = WordCloud(width = 800, height = 800,
                        background_color = 'white', collocations = False,
                        min_font_size = 10).generate(comment_words)

plot.figure(figsize = (8, 8), facecolor = None)
plot.imshow(wordcloud)
plot.title("Most common words in severe toxic comments")
plot.axis("off")
plot.tight_layout(pad = 0)

plot.show()
```













We will be evaluating model performance using the mean column-wise ROC AUC. In other words, the score is the average of the individual AUCs of each predicted column.

### 3.1 Train test split

```
In [29]: from sklearn.model_selection import train_test_split
```

```
In [31]: X_train, X_val, y_train, y_val = train_test_split(data['cleared_text1'],
                                                         data[labels], test_size=0.1, ra
X_test=test_data['cleared_text1']
tfidf = TfidfVectorizer(ngram_range = (1,3), max_features=10000)
X_train = tfidf.fit_transform(X_train)
X_val = tfidf.transform(X_val)
X_test = tfidf.transform(X_test)
feature_names = tfidf.get_feature_names()

print('Final Data dimensions after transformations:', X_train.shape, y_train.shap
```

```
Final Data dimensions after transformations: (143613, 10000) (143613, 6) (15958,
10000) (15958, 6)
```

### 3.2 Naive Bayes

#### 3.2.1 hyper parameter tuning



```

In [32]: from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import log_loss, roc_auc_score

# model = LogisticRegression(C=3)

#preds_test = np.zeros((len(test_df), len(label_col)))
alphas = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
train_roc_mean = []
val_roc_mean = []
for alp in alphas:
    print("for alpha=", alp)
    train_rocs = []
    valid_rocs = []

    preds_train = np.zeros(y_train.shape)
    preds_valid = np.zeros(y_val.shape)
    for i, label_name in enumerate(labels):
        # fit
        model = MultinomialNB(alpha=alp)
        model.fit(X_train, y_train[label_name])

        # train
        preds_train[:, i] = model.predict_proba(X_train)[:, 1]
        train_roc_class = roc_auc_score(y_train[label_name], preds_train[:, i])
        train_rocs.append(train_roc_class)

        # valid
        preds_valid[:, i] = model.predict_proba(X_val)[:, 1]
        valid_roc_class = roc_auc_score(y_val[label_name], preds_valid[:, i])
        valid_rocs.append(valid_roc_class)

    print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
    train_roc_mean.append(np.mean(train_rocs))
    print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))
    val_roc_mean.append(np.mean(valid_rocs))
    print('*'*50)

```

for alpha= 0.001

mean column-wise ROC AUC on Train data: 0.9804756171963823  
mean column-wise ROC AUC on Val data: 0.943238496064159  
\*\*\*\*\*

for alpha= 0.01

mean column-wise ROC AUC on Train data: 0.9796858736692707  
mean column-wise ROC AUC on Val data: 0.9509158804956107  
\*\*\*\*\*

for alpha= 0.1

mean column-wise ROC AUC on Train data: 0.9776564416128739  
mean column-wise ROC AUC on Val data: 0.9565998751367714  
\*\*\*\*\*

for alpha= 1

mean column-wise ROC AUC on Train data: 0.9610401339523474  
mean column-wise ROC AUC on Val data: 0.9416001762697533  
\*\*\*\*\*

for alpha= 10

mean column-wise ROC AUC on Train data: 0.8908531026672627  
mean column-wise ROC AUC on Val data: 0.8726036726128338  
\*\*\*\*\*

```
for alpha= 100
```

```
mean column-wise ROC AUC on Train data: 0.8123196668148406
```

```
mean column-wise ROC AUC on Val data: 0.8002116215150782
```

```
*****
```


```
for alpha= 1000
```

```
mean column-wise ROC AUC on Train data: 0.7781903430320467
```

```
mean column-wise ROC AUC on Val data: 0.7717305132957312
```

```
*****
```

```
In [36]: plot.plot(np.log10(alphas), train_roc_mean, label="Train AUC")
plot.plot(np.log10(alphas), val_roc_mean, label="Validation AUC")
plot.xlabel("Alpha values")
plot.ylabel("Mean ROC AUC")
plot.legend()
plot.show()
```



### 3.2.2 NB Model with tuned alpha

```

In [37]: model = MultinomialNB(0.1)

train_rocs = []
valid_rocs = []

preds_train = np.zeros(y_train.shape)
preds_valid = np.zeros(y_val.shape)
preds_test = np.zeros((len(test_data), len(labels)))

for i, label_name in enumerate(labels):
    print('\nClass:= '+label_name)
    # fit
    model.fit(X_train,y_train[label_name])

    # train
    preds_train[:,i] = model.predict_proba(X_train)[: ,1]
    train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
    print('Train ROC AUC:', train_roc_class)
    train_rocs.append(train_roc_class)

    # valid
    preds_valid[:,i] = model.predict_proba(X_val)[: ,1]
    valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
    print('Valid ROC AUC:', valid_roc_class)
    valid_rocs.append(valid_roc_class)

    # test predictions
    preds_test[:,i] = model.predict_proba(X_test)[: ,1]

print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))

```

```

Class:= toxic
Train ROC AUC: 0.9605661875124668
Valid ROC AUC: 0.9448275168382719

```

```

Class:= severe_toxic
Train ROC AUC: 0.989394185297663
Valid ROC AUC: 0.9817459891781647

```

```

Class:= obscene
Train ROC AUC: 0.9719346667954432
Valid ROC AUC: 0.9570460361743504

```

```

Class:= threat
Train ROC AUC: 0.9908676474524621
Valid ROC AUC: 0.9512356069724031

```

```

Class:= insult
Train ROC AUC: 0.9701161390044313
Valid ROC AUC: 0.9563330052424035

```

```

Class:= identity_hate
Train ROC AUC: 0.9830598236147768
Valid ROC AUC: 0.9484110964150347

```

```

mean column-wise ROC AUC on Train data: 0.9776564416128739
mean column-wise ROC AUC on Val data: 0.9565998751367714

```

### 3.2.3 Prediction on test data

```
In [60]: submission = pd.DataFrame()
submission['id'] = test_data.id
```

```
In [61]: i=0
for lab in labels:
    submission[lab] = preds_test[:,i]
    i+=1
submission.head()
```

```
Out[61]:
```

	id	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	00001cee341fdb12	0.999135	0.817937	0.997102	0.167417	0.992711	0.859244
1	0000247867823ef7	0.016884	0.000348	0.005635	0.000178	0.005653	0.000487
2	00013b17ad220c46	0.121127	0.014117	0.070228	0.001386	0.053768	0.014947
3	00017563c3f7919a	0.010603	0.000230	0.004220	0.000034	0.004207	0.000105
4	00017695ad8997eb	0.048732	0.001817	0.020911	0.000241	0.018993	0.001098

```
In [62]: submission.to_csv("test_pred_simple_NB.csv", index=False)
```

## 3.3 Logistic Regression

### 3.3.1 hyper parameter tuning

```

In [38]: alphas= [0.01,0.03, 0.1,0.3, 1,3, 10,30, 100]
train_roc_mean=[]
val_roc_mean=[]
for alp in alphas:
    print("for C=", alp)
    train_rocs = []
    valid_rocs = []

    preds_train = np.zeros(y_train.shape)
    preds_valid = np.zeros(y_val.shape)
    for i, label_name in enumerate(labels):
        # fit
        model = LogisticRegression(C=alp)
        model.fit(X_train,y_train[label_name])

        # train
        preds_train[:,i] = model.predict_proba(X_train)[:,1]
        train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
        train_rocs.append(train_roc_class)

        # valid
        preds_valid[:,i] = model.predict_proba(X_val)[:,1]
        valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
        valid_rocs.append(valid_roc_class)

    print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
    train_roc_mean.append(np.mean(train_rocs))
    print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))
    val_roc_mean.append(np.mean(valid_rocs))
    print('*'*50)

```

for C= 0.01

mean column-wise ROC AUC on Train data: 0.960654122554694  
mean column-wise ROC AUC on Val data: 0.9500028703084221  
\*\*\*\*\*

for C= 0.03

mean column-wise ROC AUC on Train data: 0.9693338936173018  
mean column-wise ROC AUC on Val data: 0.9605946857949249  
\*\*\*\*\*

for C= 0.1

mean column-wise ROC AUC on Train data: 0.9761713465601048  
mean column-wise ROC AUC on Val data: 0.9677353021834335  
\*\*\*\*\*

for C= 0.3

mean column-wise ROC AUC on Train data: 0.9817582166953253  
mean column-wise ROC AUC on Val data: 0.9718489365831274  
\*\*\*\*\*

for C= 1

mean column-wise ROC AUC on Train data: 0.9869967581598035  
mean column-wise ROC AUC on Val data: 0.9736092438933243  
\*\*\*\*\*

for C= 3

mean column-wise ROC AUC on Train data: 0.9906705683605713  
mean column-wise ROC AUC on Val data: 0.9719974527003972  
\*\*\*\*\*

for C= 10

mean column-wise ROC AUC on Train data: 0.9930931741190242

mean column-wise ROC AUC on Val data: 0.9671086961125234

\*\*\*\*\*

for C= 30

mean column-wise ROC AUC on Train data: 0.9941810941135795

mean column-wise ROC AUC on Val data: 0.9606485768011757

\*\*\*\*\*


for C= 100

mean column-wise ROC AUC on Train data: 0.9948237638091584

mean column-wise ROC AUC on Val data: 0.9529761404682379

\*\*\*\*\*

```
In [39]: plot.plot(np.log10(alphas), train_roc_mean, label="Train AUC")
plot.plot(np.log10(alphas), val_roc_mean, label="Validation AUC")
plot.xlabel("C values")
plot.ylabel("Mean ROC AUC")
plot.legend()
plot.show()
```



```

In [40]: alphas= [0.5, 1, 1.5, 2, 2.5, 3]
train_roc_mean=[]
val_roc_mean=[]
for alp in alphas:
    print("for C=", alp)
    train_rocs = []
    valid_rocs = []

    preds_train = np.zeros(y_train.shape)
    preds_valid = np.zeros(y_val.shape)
    for i, label_name in enumerate(labels):
        # fit
        model = LogisticRegression(C=alp)
        model.fit(X_train,y_train[label_name])

        # train
        preds_train[:,i] = model.predict_proba(X_train)[:,1]
        train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
        train_rocs.append(train_roc_class)

        # valid
        preds_valid[:,i] = model.predict_proba(X_val)[:,1]
        valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
        valid_rocs.append(valid_roc_class)

    print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
    train_roc_mean.append(np.mean(train_rocs))
    print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))
    val_roc_mean.append(np.mean(valid_rocs))
    print('*'*50)

```

for C= 0.5

```

mean column-wise ROC AUC on Train data:  0.9841146113528881
mean column-wise ROC AUC on Val data: 0.9730250039538642
*****

```

for C= 1

```

mean column-wise ROC AUC on Train data:  0.9869967581598035
mean column-wise ROC AUC on Val data: 0.9736092438933243
*****

```

for C= 1.5

```

mean column-wise ROC AUC on Train data:  0.988499883342388
mean column-wise ROC AUC on Val data: 0.9733812943884753
*****

```

for C= 2

```

mean column-wise ROC AUC on Train data:  0.9894664565906254
mean column-wise ROC AUC on Val data: 0.9729600074733468
*****

```

for C= 2.5

```

mean column-wise ROC AUC on Train data:  0.9901516330301168
mean column-wise ROC AUC on Val data: 0.9724813103058153
*****

```

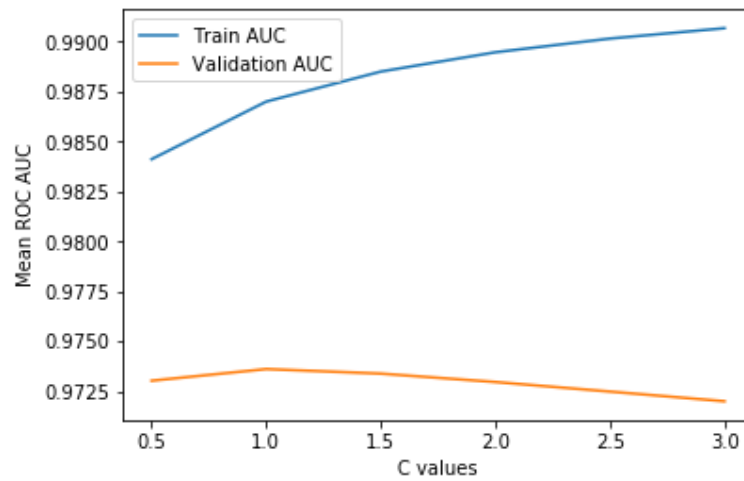
for C= 3

```

mean column-wise ROC AUC on Train data:  0.9906705683605713
mean column-wise ROC AUC on Val data: 0.9719974527003972
*****

```

```
In [41]: plot.plot(alphas, train_roc_mean, label="Train AUC")
plot.plot(alphas, val_roc_mean, label="Validation AUC")
plot.xlabel("C values")
plot.ylabel("Mean ROC AUC")
plot.legend()
plot.show()
```



### 3.3.2 LR Model with tuned C



```

In [67]: model = LogisticRegression(C=1)

train_rocs = []
valid_rocs = []

preds_train = np.zeros(y_train.shape)
preds_valid = np.zeros(y_val.shape)
preds_test = np.zeros((len(test_data), len(labels)))

for i, label_name in enumerate(labels):
    print('\nClass:= '+label_name)
    # fit
    model.fit(X_train,y_train[label_name])

    # train
    preds_train[:,i] = model.predict_proba(X_train)[:,1]
    train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
    print('Train ROC AUC:', train_roc_class)
    train_rocs.append(train_roc_class)

    # valid
    preds_valid[:,i] = model.predict_proba(X_val)[:,1]
    valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
    print('Valid ROC AUC:', valid_roc_class)
    valid_rocs.append(valid_roc_class)

    # test predictions
    preds_test[:,i] = model.predict_proba(X_test)[:,1]

print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))

```

```

Class:= toxic
Train ROC AUC: 0.9780188873089515
Valid ROC AUC: 0.9648575882452063

```

```

Class:= severe_toxic
Train ROC AUC: 0.9905106590102156
Valid ROC AUC: 0.9814785081804605

```

```

Class:= obscene
Train ROC AUC: 0.9901677297191414
Valid ROC AUC: 0.9794587933285638

```

```

Class:= threat
Train ROC AUC: 0.9938611663006507
Valid ROC AUC: 0.9805211072863553

```

```

Class:= insult
Train ROC AUC: 0.9830649218117186
Valid ROC AUC: 0.9730099791904906

```

```

Class:= identity_hate
Train ROC AUC: 0.9863571848081442
Valid ROC AUC: 0.9623294871288692

```

```

mean column-wise ROC AUC on Train data: 0.9869967581598035
mean column-wise ROC AUC on Val data: 0.9736092438933243

```

### 3.3.3 Prediction on test data

```
In [68]: submission = pd.DataFrame()
submission['id'] = test_data.id
```

```
In [69]: i=0
for lab in labels:
    submission[lab]= preds_test[:,i]
    i+=1
submission.head()
```

```
Out[69]:
```

	id	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	00001cee341fdb12	0.999577	0.290275	0.998708	0.055111	0.977366	0.517051
1	0000247867823ef7	0.009613	0.003850	0.005704	0.001660	0.009079	0.003409
2	00013b17ad220c46	0.119847	0.007554	0.028889	0.002276	0.026148	0.008345
3	00017563c3f7919a	0.004832	0.002303	0.003598	0.001123	0.004191	0.000947
4	00017695ad8997eb	0.038326	0.001228	0.006754	0.000488	0.011149	0.000749

```
In [70]: submission.to_csv("test_pred_LR1.csv", index=False)
```

## 3.4 SVM

### 3.4.1 hyper parameter tuning

```
In [42]: from sklearn.linear_model import SGDClassifier
from sklearn.calibration import CalibratedClassifierCV
import warnings
warnings.filterwarnings("ignore")

alphas = [0.000001, 0.00001, 0.0001, 0.001, 0.01, 0.1, 1, 10]
train_roc_mean = []
val_roc_mean = []
for alp in alphas:
    print("for alpha=", alp)
    train_rocs = []
    valid_rocs = []

    preds_train = np.zeros(y_train.shape)
    preds_valid = np.zeros(y_val.shape)
    for i, label_name in enumerate(labels):
        # fit
        modelsgd = SGDClassifier(alpha=alp)
        modelsgd.fit(X_train, y_train[label_name])
        model = CalibratedClassifierCV(modelsgd, cv='prefit', method='sigmoid')
        model.fit(X_train, y_train[label_name])

        # train
        preds_train[:, i] = model.predict_proba(X_train)[:, 1]
        train_roc_class = roc_auc_score(y_train[label_name], preds_train[:, i])
        train_rocs.append(train_roc_class)

        # valid
        preds_valid[:, i] = model.predict_proba(X_val)[:, 1]
        valid_roc_class = roc_auc_score(y_val[label_name], preds_valid[:, i])
        valid_rocs.append(valid_roc_class)

    print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
    train_roc_mean.append(np.mean(train_rocs))
    print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))
    val_roc_mean.append(np.mean(valid_rocs))
    print('*'*50)
```

for alpha= 1e-06

mean column-wise ROC AUC on Train data: 0.9851218556403406  
mean column-wise ROC AUC on Val data: 0.9473076888304311  
\*\*\*\*\*

for alpha= 1e-05

mean column-wise ROC AUC on Train data: 0.9859174036214874  
mean column-wise ROC AUC on Val data: 0.9575892729460976  
\*\*\*\*\*

for alpha= 0.0001

mean column-wise ROC AUC on Train data: 0.9832336949426733  
mean column-wise ROC AUC on Val data: 0.9618978674867517  
\*\*\*\*\*

for alpha= 0.001

mean column-wise ROC AUC on Train data: 0.9808165832286756  
mean column-wise ROC AUC on Val data: 0.961385832607372  
\*\*\*\*\*

for alpha= 0.01

mean column-wise ROC AUC on Train data: 0.9806687375183228  
mean column-wise ROC AUC on Val data: 0.9586291390180869  
\*\*\*\*\*

for alpha= 0.1

```

mean column-wise ROC AUC on Train data: 0.7620480932201367
mean column-wise ROC AUC on Val data: 0.7538399460948709
*****
for alpha= 1

mean column-wise ROC AUC on Train data: 0.7606578909608733
mean column-wise ROC AUC on Val data: 0.7587358482057341
*****
for alpha= 10

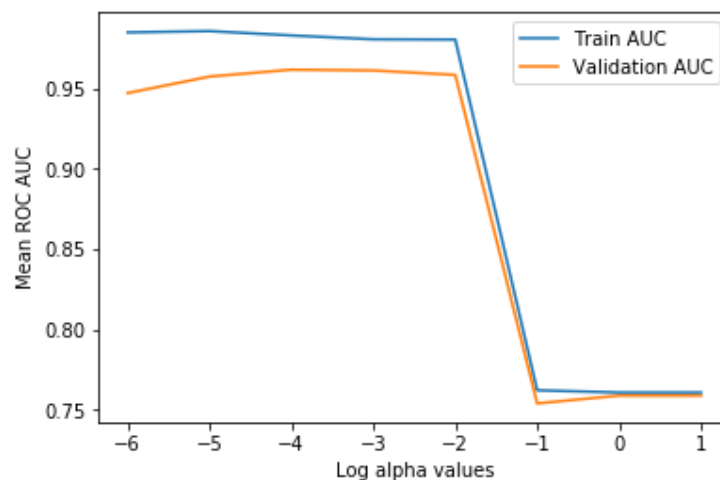
mean column-wise ROC AUC on Train data: 0.7606578909608733
mean column-wise ROC AUC on Val data: 0.7587358482057341
*****

```

```

In [44]: plot.plot(np.log10(alphas), train_roc_mean, label="Train AUC")
plot.plot(np.log10(alphas), val_roc_mean, label="Validation AUC")
plot.xlabel("Log alpha values")
plot.ylabel("Mean ROC AUC")
plot.legend()
plot.show()

```



### 3.4.2 SVM Model with tuned alpha

```

In [83]: modelsgd = SGDClassifier(alpha=0.001)

train_rocs = []
valid_rocs = []

preds_train = np.zeros(y_train.shape)
preds_valid = np.zeros(y_val.shape)
preds_test = np.zeros((len(test_data), len(labels)))

for i, label_name in enumerate(labels):
    print('\nClass:= '+label_name)
    # fit
    modelsgd.fit(X_train,y_train[label_name])
    model=CalibratedClassifierCV(modelsgd, cv='prefit', method= 'sigmoid')
    model.fit(X_train,y_train[label_name])

    # train
    preds_train[:,i] = model.predict_proba(X_train)[:,1]
    train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
    print('Train ROC AUC:', train_roc_class)
    train_rocs.append(train_roc_class)

    # valid
    preds_valid[:,i] = model.predict_proba(X_val)[:,1]
    valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
    print('Valid ROC AUC:', valid_roc_class)
    valid_rocs.append(valid_roc_class)

    # test predictions
    preds_test[:,i] = model.predict_proba(X_test)[:,1]

print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))

```

```

Class:= toxic
Train ROC AUC: 0.964212852756118
Valid ROC AUC: 0.9593357901042857

```

```

Class:= severe_toxic
Train ROC AUC: 0.9886401388353874
Valid ROC AUC: 0.9639282022490694

```

```

Class:= obscene
Train ROC AUC: 0.9839463445475733
Valid ROC AUC: 0.9744046772538126

```

```

Class:= threat
Train ROC AUC: 0.997208778597442
Valid ROC AUC: 0.9669374792786244

```

```

Class:= insult
Train ROC AUC: 0.9673047244432166
Valid ROC AUC: 0.9617462981927624

```

```

Class:= identity_hate
Train ROC AUC: 0.9874475795611554
Valid ROC AUC: 0.9447501283164876

```

```

mean column-wise ROC AUC on Train data: 0.9814600697901489
mean column-wise ROC AUC on Val data: 0.9618504292325071

```

### 3.4.3 Prediction on test data

```
In [84]: submission = pd.DataFrame()
submission['id'] = test_data.id
```

```
In [85]: i=0
for lab in labels:
    submission[lab] = preds_test[:,i]
    i+=1
submission.head()
```

```
Out[85]:
```

	id	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	00001cee341fdb12	1.000000	0.346314	0.998408	0.026987	0.768006	0.074627
1	0000247867823ef7	0.031744	0.000923	0.017322	0.000014	0.028347	0.004650
2	00013b17ad220c46	0.077417	0.002253	0.023325	0.000027	0.032195	0.012066
3	00017563c3f7919a	0.022035	0.004957	0.015411	0.000081	0.023885	0.002916
4	00017695ad8997eb	0.046499	0.001832	0.017734	0.000030	0.022184	0.004450

```
In [86]: submission.to_csv("test_pred_SVM.csv", index=False)
```

## 3.5 Logistic Regression2

### 3.5.1 hyper parameter tuning

```
In [45]: X_train, X_val, y_train, y_val = train_test_split(data['cleared_text1'],
                                                         data[labels], test_size=0.1, ra
X_test=test_data['cleared_text1']

vec = TfidfVectorizer(ngram_range=(1,3),
                     min_df=10, max_df=0.9, strip_accents='unicode', use_idf=1,
                     smooth_idf=1, sublinear_tf=1)
X_train = vec.fit_transform(X_train)
X_val = vec.transform(X_val)
X_test = vec.transform(X_test)
```

```

In [46]: alphas= [0.1,0.5, 1,2,3,4]
train_roc_mean=[]
val_roc_mean=[]
for alp in alphas:
    print("for C=", alp)
    train_rocs = []
    valid_rocs = []

    preds_train = np.zeros(y_train.shape)
    preds_valid = np.zeros(y_val.shape)
    for i, label_name in enumerate(labels):
        # fit
        model = LogisticRegression(C=alp)
        model.fit(X_train,y_train[label_name])

        # train
        preds_train[:,i] = model.predict_proba(X_train)[:,1]
        train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
        train_rocs.append(train_roc_class)

        # valid
        preds_valid[:,i] = model.predict_proba(X_val)[:,1]
        valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
        valid_rocs.append(valid_roc_class)

    print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
    train_roc_mean.append(np.mean(train_rocs))
    print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))
    val_roc_mean.append(np.mean(valid_rocs))
    print('*'*50)

```

for C= 0.1

```

mean column-wise ROC AUC on Train data:  0.979391409671884
mean column-wise ROC AUC on Val data: 0.9708983570784514
*****

```

for C= 0.5

```

mean column-wise ROC AUC on Train data:  0.9876736643175636
mean column-wise ROC AUC on Val data: 0.976517399661868
*****

```

for C= 1

```

mean column-wise ROC AUC on Train data:  0.9912894855622932
mean column-wise ROC AUC on Val data: 0.9779113180471359
*****

```

for C= 2

```

mean column-wise ROC AUC on Train data:  0.9944158298533184
mean column-wise ROC AUC on Val data: 0.9781858189624962
*****

```

for C= 3

```

mean column-wise ROC AUC on Train data:  0.995873858798145
mean column-wise ROC AUC on Val data: 0.9778106754208529
*****

```

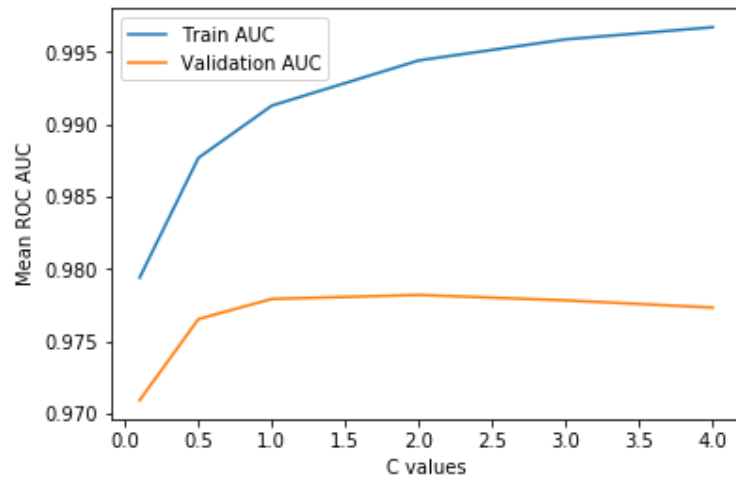
for C= 4

```

mean column-wise ROC AUC on Train data:  0.9967144796589809
mean column-wise ROC AUC on Val data: 0.977311156910332
*****

```

```
In [47]: plot.plot(alphas, train_roc_mean, label="Train AUC")
plot.plot(alphas, val_roc_mean, label="Validation AUC")
plot.xlabel("C values")
plot.ylabel("Mean ROC AUC")
plot.legend()
plot.show()
```



### 3.5.2 LR2 Model with tuned C



```

In [89]: model = LogisticRegression(C=2)

train_rocs = []
valid_rocs = []

preds_train = np.zeros(y_train.shape)
preds_valid = np.zeros(y_val.shape)
preds_test = np.zeros((len(test_data), len(labels)))

for i, label_name in enumerate(labels):
    print('\nClass:= '+label_name)
    # fit
    model.fit(X_train,y_train[label_name])

    # train
    preds_train[:,i] = model.predict_proba(X_train)[:,1]
    train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
    print('Train ROC AUC:', train_roc_class)
    train_rocs.append(train_roc_class)

    # valid
    preds_valid[:,i] = model.predict_proba(X_val)[:,1]
    valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
    print('Valid ROC AUC:', valid_roc_class)
    valid_rocs.append(valid_roc_class)

    # test predictions
    preds_test[:,i] = model.predict_proba(X_test)[:,1]

print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))

```

```

Class:= toxic
Train ROC AUC: 0.9914819516826907
Valid ROC AUC: 0.9693451357265814

```

```

Class:= severe_toxic
Train ROC AUC: 0.9944824878091023
Valid ROC AUC: 0.9872948267503289

```

```

Class:= obscene
Train ROC AUC: 0.9961164151823908
Valid ROC AUC: 0.9831184030941582

```

```

Class:= threat
Train ROC AUC: 0.9979666252465946
Valid ROC AUC: 0.9868769062333096

```

```

Class:= insult
Train ROC AUC: 0.9922503968306162
Valid ROC AUC: 0.9743687567842443

```

```

Class:= identity_hate
Train ROC AUC: 0.9941971023685163
Valid ROC AUC: 0.968110885186355

```

```

mean column-wise ROC AUC on Train data: 0.9944158298533184
mean column-wise ROC AUC on Val data: 0.9781858189624962

```

### 3.5.3 Prediction on test data

```
In [90]: submission = pd.DataFrame()
submission['id'] = test_data.id
```

```
In [91]: i=0
for lab in labels:
    submission[lab]= preds_test[:,i]
    i+=1
submission.head()
```

```
Out[91]:
```

	id	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	00001cee341fdb12	0.999300	0.210607	0.997974	0.038784	0.949908	0.364595
1	0000247867823ef7	0.003894	0.002484	0.003067	0.001094	0.006326	0.002330
2	00013b17ad220c46	0.119422	0.005506	0.022499	0.001670	0.019493	0.007020
3	00017563c3f7919a	0.005042	0.001977	0.002873	0.001180	0.003152	0.000870
4	00017695ad8997eb	0.017262	0.001570	0.005401	0.000616	0.010236	0.001010

```
In [92]: submission.to_csv("test_pred_LR2.csv", index=False)
```

## 3.6 NB-SVM

### 3.6.1 NB-SVM modeling

```
In [133]: # ref: https://www.kaggle.com/jhoward/nb-svm-strong-linear-baseline
# NB-SVM
def pr(y_i, y):
    p = x[y==y_i].sum(0)
    return (p+1) / ((y==y_i).sum()+1)
```

```
In [134]: x = X_train
test_x = X_val
```

```
In [135]: def get_md1(y):
    y = y.values
    r = np.log(pr(1,y) / pr(0,y))
    m = LogisticRegression(C=4, dual=True)
    x_nb = x.multiply(r)
    return m.fit(x_nb, y), r
```

```

In [136]: train_rocs = []
          valid_rocs = []

          preds_train = np.zeros(y_train.shape)
          preds_valid = np.zeros(y_val.shape)
          preds_test = np.zeros((len(test_data), len(labels)))

          for i, label_name in enumerate(labels):
              print('\nClass:= '+label_name)
              # fit
              model,r = get_mdl(y_train[label_name])
              #model.fit(X_train,)

              # train
              preds_train[:,i] = model.predict_proba(X_train.multiply(r))[:,1]
              train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
              print('Train ROC AUC:', train_roc_class)
              train_rocs.append(train_roc_class)

              # valid
              preds_valid[:,i] = model.predict_proba(X_val.multiply(r))[:,1]
              valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
              print('Valid ROC AUC:', valid_roc_class)
              valid_rocs.append(valid_roc_class)

              # test predictions
              preds_test[:,i] = model.predict_proba(X_test.multiply(r))[:,1]

          print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
          print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))

```

```

Class:= toxic
Train ROC AUC: 0.9793015297520724
Valid ROC AUC: 0.9662317652500112

```

```

Class:= severe_toxic
Train ROC AUC: 0.9940141622446979
Valid ROC AUC: 0.9804576920290662

```

```

Class:= obscene
Train ROC AUC: 0.990485960236388
Valid ROC AUC: 0.9804965234460078

```

```

Class:= threat
Train ROC AUC: 0.9993968744373798
Valid ROC AUC: 0.9705897716624327

```

```

Class:= insult
Train ROC AUC: 0.9835116969658667
Valid ROC AUC: 0.9720578933244072

```

```

Class:= identity_hate
Train ROC AUC: 0.991683789323279
Valid ROC AUC: 0.9555440125552748

```

```

mean column-wise ROC AUC on Train data: 0.9897323354932807
mean column-wise ROC AUC on Val data: 0.9708962763778666

```

### 3.6.2 Prediction on test data

```
In [137]: submission = pd.DataFrame()
submission['id'] = test_data.id
```

```
In [138]: i=0
for lab in labels:
    submission[lab] = preds_test[:,i]
    i+=1
submission.head()
```

```
Out[138]:
```

	id	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	00001cee341fdb12	0.999996	0.435986	0.999976	0.056048	0.992826	0.783812
1	0000247867823ef7	0.002547	0.001184	0.001263	0.000135	0.004204	0.000691
2	00013b17ad220c46	0.047265	0.001802	0.010657	0.000122	0.013883	0.006360
3	00017563c3f7919a	0.001564	0.000909	0.001230	0.000177	0.002566	0.000092
4	00017695ad8997eb	0.038281	0.000446	0.004149	0.000029	0.011418	0.000120

```
In [139]: submission.to_csv("test_pred_NB-SVM.csv", index=False)
```

## 3.7 Random Forest

### 3.7.1 Random Forest Modeling

```
In [105]: X_train, X_val, y_train, y_val = train_test_split(data['cleared_text1'],
                                                         data[labels], test_size=0.1, ra
X_test=test_data['cleared_text1']
tfidf = TfidfVectorizer(ngram_range = (1,3), min_df = 150)
X_train = tfidf.fit_transform(X_train)
X_val = tfidf.transform(X_val)
X_test = tfidf.transform(X_test)

print('Final Data dimensions after transformations:', X_train.shape, y_train.shap
```

Final Data dimensions after transformations: (143613, 5075) (143613, 6) (15958, 5075) (15958, 6)

```
In [106]: from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

In [107]:

```
model=RandomForestClassifier(n_estimators=150,max_depth=60,max_features='auto')

train_rocs = []
valid_rocs = []

preds_train = np.zeros(y_train.shape)
preds_valid = np.zeros(y_val.shape)
preds_test = np.zeros((len(test_data), len(labels)))

for i, label_name in enumerate(labels):
    print('\nClass:= '+label_name)
    # fit
    model.fit(X_train,y_train[label_name])

    # train
    preds_train[:,i] = model.predict_proba(X_train)[:,1]
    train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
    print('Train ROC AUC:', train_roc_class)
    train_rocs.append(train_roc_class)

    # valid
    preds_valid[:,i] = model.predict_proba(X_val)[:,1]
    valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
    print('Valid ROC AUC:', valid_roc_class)
    valid_rocs.append(valid_roc_class)

    # test predictions
    preds_test[:,i] = model.predict_proba(X_test)[:,1]

print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))
```

```
Class:= toxic
Train ROC AUC: 0.9922514391113061
Valid ROC AUC: 0.9443152751214732
```

```
Class:= severe_toxic
Train ROC AUC: 0.9973162854864863
Valid ROC AUC: 0.9773060832423155
```

```
Class:= obscene
Train ROC AUC: 0.9971215749522314
Valid ROC AUC: 0.9712960792939039
```

```
Class:= threat
Train ROC AUC: 0.9993902075802632
Valid ROC AUC: 0.9324656384955455
```

```
Class:= insult
Train ROC AUC: 0.994888394983711
Valid ROC AUC: 0.9598697093263683
```

```
Class:= identity_hate
Train ROC AUC: 0.9975863111512552
Valid ROC AUC: 0.9343986891977258
```

```
mean column-wise ROC AUC on Train data: 0.9964257022108756
mean column-wise ROC AUC on Val data: 0.9532752457795555
```

### 3.7.2 Prediction on test data

```
In [108]: submission = pd.DataFrame()
submission['id'] = test_data.id
```

```
In [109]: i=0
for lab in labels:
    submission[lab] = preds_test[:,i]
    i+=1
submission.head()
```

```
Out[109]:
```

	id	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	00001cee341fdb12	0.877610	0.172997	0.809736	0.074539	0.697454	0.108143
1	0000247867823ef7	0.072948	0.003768	0.031091	0.000652	0.030761	0.003552
2	00013b17ad220c46	0.078572	0.003775	0.032298	0.000652	0.031526	0.003552
3	00017563c3f7919a	0.031431	0.002003	0.023172	0.000631	0.012816	0.002525
4	00017695ad8997eb	0.040931	0.008226	0.012516	0.000460	0.014427	0.001736

```
In [110]: submission.to_csv("test_pred_NB-SVM.csv", index=False)
```

## 3.8 LR WordGrams and CharacterGrams

### 3.8.1 featurization

```
In [30]: X_train, X_val, y_train, y_val = train_test_split(data['cleared_text1'],
                                                         data[labels], test_size=0.1, ra
X_test=test_data['cleared_text1']

word_vec = TfidfVectorizer(ngram_range=(1,2), analyzer='word',max_features=12000,
                           strip_accents='unicode', use_idf=1,
                           smooth_idf=1, sublinear_tf=1 )
X_trainw = word_vec.fit_transform(X_train)
X_valw = word_vec.transform(X_val)
X_testw = word_vec.transform(X_test)
```

```
In [31]: char_vec = TfidfVectorizer(ngram_range=(2,6), analyzer='char',max_features=50000,
                                    strip_accents='unicode', use_idf=1,
                                    smooth_idf=1, sublinear_tf=1 )
X_trainc = char_vec.fit_transform(X_train)
X_valc = char_vec.transform(X_val)
X_testc = char_vec.transform(X_test)
```

```
In [32]: from scipy.sparse import hstack

X_train=hstack([X_trainw, X_trainc])
X_val=hstack([X_valw, X_valc])
X_test= hstack([X_testw, X_testc])

print(X_train.shape)
```

```
(143613, 62000)
```

### 3.8.2 LR Model

```
In [35]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score
alphas= [1,2,3]
train_roc_mean=[]
val_roc_mean=[]
for alp in alphas:
    print("for C=", alp)
    train_rocs = []
    valid_rocs = []

    preds_train = np.zeros(y_train.shape)
    preds_valid = np.zeros(y_val.shape)
    for i, label_name in enumerate(labels):
        # fit
        model = LogisticRegression(C=alp)
        model.fit(X_train,y_train[label_name])

        # train
        preds_train[:,i] = model.predict_proba(X_train)[:,1]
        train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
        train_rocs.append(train_roc_class)

        # valid
        preds_valid[:,i] = model.predict_proba(X_val)[:,1]
        valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
        valid_rocs.append(valid_roc_class)

    print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
    train_roc_mean.append(np.mean(train_rocs))
    print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))
    val_roc_mean.append(np.mean(valid_rocs))
    print('*'*50)
```

for C= 1

```
mean column-wise ROC AUC on Train data:  0.9940603352616578
mean column-wise ROC AUC on Val data: 0.9816312508927254
*****
```


for C= 2

```
mean column-wise ROC AUC on Train data:  0.9962599440866207
mean column-wise ROC AUC on Val data: 0.9812192767142355
*****
```

for C= 3

```
mean column-wise ROC AUC on Train data:  0.9972877396027896
mean column-wise ROC AUC on Val data: 0.9805176126183208
*****
```

```
In [36]: plot.plot(alphas, train_roc_mean, label="Train AUC")
plot.plot(alphas, val_roc_mean, label="Validation AUC")
plot.xlabel("C values")
plot.ylabel("Mean ROC AUC")
plot.legend()
plot.show()
```





```

In [124]: model = LogisticRegression(C=1)

train_rocs = []
valid_rocs = []

preds_train = np.zeros(y_train.shape)
preds_valid = np.zeros(y_val.shape)
preds_test = np.zeros((len(test_data), len(labels)))

for i, label_name in enumerate(labels):
    print('\nClass:= '+label_name)
    # fit
    model.fit(X_train,y_train[label_name])

    # train
    preds_train[:,i] = model.predict_proba(X_train)[:,1]
    train_roc_class = roc_auc_score(y_train[label_name],preds_train[:,i])
    print('Train ROC AUC:', train_roc_class)
    train_rocs.append(train_roc_class)

    # valid
    preds_valid[:,i] = model.predict_proba(X_val)[:,1]
    valid_roc_class = roc_auc_score(y_val[label_name],preds_valid[:,i])
    print('Valid ROC AUC:', valid_roc_class)
    valid_rocs.append(valid_roc_class)

    # test predictions
    preds_test[:,i] = model.predict_proba(X_test)[:,1]

print('\nmean column-wise ROC AUC on Train data: ', np.mean(train_rocs))
print('mean column-wise ROC AUC on Val data:', np.mean(valid_rocs))

```

```

Class:= toxic
Train ROC AUC: 0.989207086204653
Valid ROC AUC: 0.9724931171299626

```

```

Class:= severe_toxic
Train ROC AUC: 0.9947881038338992
Valid ROC AUC: 0.9876049723602969

```

```

Class:= obscene
Train ROC AUC: 0.9964868315941662
Valid ROC AUC: 0.9893221606197531

```

```

Class:= threat
Train ROC AUC: 0.997415809511625
Valid ROC AUC: 0.9836563647004768

```

```

Class:= insult
Train ROC AUC: 0.9912788410184878
Valid ROC AUC: 0.9810145268377137

```

```

Class:= identity_hate
Train ROC AUC: 0.9951853394071156
Valid ROC AUC: 0.9756963637081492

```

```

mean column-wise ROC AUC on Train data: 0.9940603352616578
mean column-wise ROC AUC on Val data: 0.9816312508927254

```

### 3.8.3 Prediction on test data

```
In [125]: submission = pd.DataFrame()
submission['id'] = test_data.id
```

```
In [126]: i=0
for lab in labels:
    submission[lab] = preds_test[:,i]
    i+=1
submission.head()
```

```
Out[126]:
```

	id	toxic	severe_toxic	obscene	threat	insult	identity_hate
0	00001cee341fdb12	0.999955	0.279281	0.999892	0.064325	0.991839	0.439833
1	0000247867823ef7	0.008825	0.003128	0.003799	0.001344	0.007943	0.002206
2	00013b17ad220c46	0.051237	0.003590	0.011305	0.001422	0.007983	0.003532
3	00017563c3f7919a	0.003036	0.001953	0.002554	0.000785	0.002516	0.000345
4	00017695ad8997eb	0.015663	0.001920	0.004481	0.000605	0.006929	0.001557

```
In [127]: submission.to_csv("test_pred_LR_W&C-Grams.csv", index=False)
```

## 4. Comparing Model Result

```
In [40]: from prettytable import PrettyTable

x = PrettyTable()

x.field_names = ["Model Description/(Score- Mean ROC AUC)", "Kaggle Private Score", "Kaggle Public Score", "Local Val Score"]

x.add_row(["Naive Bayes", 0.95570, 0.95347, 0.9565])
x.add_row(["SVM", 0.95643, 0.96053, 0.9619])
x.add_row(["Naive Bayes - SVM", 0.96727, 0.96630, 0.9709])
x.add_row(["Logistic Regression Limited features (10K)", 0.97121, 0.97055, 0.9736])
x.add_row(["Logistic Regression v2", 0.97489, 0.97396, 0.9782])
x.add_row(["Logistic Regression Word and Character Grams", 0.97855, 0.97706, 0.9816])
x.add_row(["Random Forest", 0.95350, 0.95273, 0.9533])

print(x)
```

Model Description/(Score- Mean ROC AUC)	Kaggle Private Score	Kaggle Public Score	Local Val Score
Naive Bayes	0.9557	0.95347	0.9565
SVM	0.95643	0.96053	0.9619
Naive Bayes - SVM	0.96727	0.96630	0.9709
Logistic Regression Limited features (10K)	0.97121	0.97055	0.9736
Logistic Regression v2	0.97489	0.97396	0.9782
Logistic Regression Word and Character Grams	0.97855	0.97706	0.9816
Random Forest	0.9535	0.95273	0.9533

## References

- <https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge/> (<https://www.kaggle.com/c/jigsaw-toxic-comment-classification-challenge/>)
- <https://medium.com/@datamonsters/text-preprocessing-in-python-steps-tools-and-examples-bf025f872908> (<https://medium.com/@datamonsters/text-preprocessing-in-python-steps-tools-and-examples-bf025f872908>)
- <https://www.kaggle.com/asrsaiteja/toxic-comments-featureeng-eda-with-nb-baseline> (<https://www.kaggle.com/asrsaiteja/toxic-comments-featureeng-eda-with-nb-baseline>)
- <https://sijunhe.github.io/blog/2018/04/03/nb-svm/> (<https://sijunhe.github.io/blog/2018/04/03/nb-svm/>)
- <https://www.kaggle.com/jhoward/nb-svm-strong-linear-baseline> (<https://www.kaggle.com/jhoward/nb-svm-strong-linear-baseline>)

In [ ]:

